

# Kashyap Chitta

Postdoctoral Researcher

✉ [kchitta@nvidia.com](mailto:kchitta@nvidia.com)    <https://kashyap7x.github.io>    📍 Tübingen, Germany, 72076

## Employment

- |            |  |
|------------|--|
| 2025 – Now | <b>NVIDIA, Germany</b><br><i>Postdoctoral Researcher; <a href="#">Autonomous Vehicle Research Group</a></i><br><i>Visiting Researcher; <a href="#">Autonomous Vision Group</a>, <a href="#">University of Tübingen</a></i> <ul style="list-style-type: none"><li>• <i>Role:</i> Research focused on simulation-based training and evaluation of embodied AI systems.</li></ul> |
|------------|--|

## Education

- |             |   |
|-------------|---|
| 2019 – 2025 | <b>University of Tübingen, Germany</b><br><i>PhD in Computer Science; <a href="#">Autonomous Vision Group</a></i> <ul style="list-style-type: none"><li>• <i>Advisor:</i> <a href="#">Prof. Andreas Geiger</a></li><li>• <i>Scholarship:</i> <a href="#">International Max Planck Research School for Intelligent Systems (IMPRS-IS)</a></li><li>• <i>Thesis:</i> Scalability-Driven Design for Autonomous Vehicles</li></ul>   |
| 2017 – 2018 | <b>Carnegie Mellon University, USA</b><br><i>Master of Science in Computer Vision</i> <ul style="list-style-type: none"><li>• <i>Advisor:</i> <a href="#">Prof. Martial Hebert</a></li><li>• <i>Thesis project:</i> Exploiting Synthetic Data for Street Scene Segmentation</li><li>• <i>GPA:</i> 4.15/4.33</li><li>• <i>Selected courses:</i> Visual Learning and Recognition, Deep Reinforcement Learning, Geometry Based Methods in Vision, Statistical Techniques in Robotics</li></ul> |
| 2013 – 2017 | <b>RV College of Engineering, India</b><br><i>Bachelor of Engineering in Electronics and Communication</i> <ul style="list-style-type: none"><li>• <i>Thesis project:</i> Monocular Visual SLAM with a Rotating Mirror</li><li>• <i>GPA:</i> 9.11/10.0</li></ul>  |

## Awards

- |      |   |
|------|---|
| 2025 | <ul style="list-style-type: none"><li>• Our approach DiffusionLTF ranked <b>second</b> on the <a href="#">2025 Waymo Vision-based End-to-End Driving Challenge</a>.</li><li>• Our approach SHRED ranked <b>third</b> on the <a href="#">2025 Waymo Scenario Generation Challenge</a>.</li></ul>   |
| 2024 | <ul style="list-style-type: none"><li>• I was named an <a href="#">outstanding reviewer</a> at ECCV 2024 (198/7293 reviewers, top 3%).</li><li>• Our approach GenDM ranked <b>second</b> on the <a href="#">2024 Dataset Distillation Challenge</a> generative track and won the <b>best paper award</b> at the challenge's ECCV workshop.</li><li>• Our approach TF++ ranked <b>first</b> on the <a href="#">2024 CARLA AD Challenge</a> map track (40 participating teams).</li></ul> |

2023	<ul style="list-style-type: none"> <li>• Our approach TF++ ranked <b>second</b> on the <a href="#">2023 CARLA AD Challenge</a> (20 participating teams).</li> <li>• I was named a <a href="#">top reviewer</a> at NeurIPS 2023 (1196/11725 reviewers, top 10%).</li> <li>• I was named an <a href="#">outstanding reviewer</a> at ICCV 2023 (130/7000 reviewers, top 2%).</li> <li>• I was selected for the <a href="#">doctoral consortium</a> at ICCV 2023 (38 accepted participants).</li> <li>• Our approach PDM ranked <b>first</b> on the <a href="#">2023 nuPlan Planning Challenge</a> (52 participating teams).</li> <li>• I was named an <a href="#">outstanding reviewer</a> at CVPR 2023 (232/7000 reviewers, top 3%).</li> <li>• I was selected as a <a href="#">2023 RSS Pioneer</a>, (30/135 applicants, 22% acceptance rate).</li> </ul>
2022	<ul style="list-style-type: none"> <li>• Our approach MapTF++ ranked <b>first</b> on the <a href="#">2022 CARLA AD Challenge</a> map track.</li> </ul>
2021	<ul style="list-style-type: none"> <li>• Our approach TransFuser ranked <b>second</b> on the <a href="#">2021 CARLA AD Challenge</a>, (100+ participating teams).</li> <li>• Our new <a href="#">computer vision lecture</a> won the 2021 CS teaching award at the University of Tübingen.</li> </ul>
2020	<ul style="list-style-type: none"> <li>• Our approach NEAT ranked <b>second</b> on the <a href="#">2020 CARLA AD Challenge</a> (45 participating teams).</li> </ul>

## Internships

---

Jan 2019 – Aug 2019	<b>NVIDIA, Santa Clara, USA</b> <i>Deep Learning Intern; <a href="#">Autonomous Vehicle Applied Research</a></i> <ul style="list-style-type: none"> <li>• <i>Mentor:</i> <a href="#">Dr. José M. Álvarez</a></li> <li>• <i>Role:</i> Research and development of an automatic dataset curation engine for the internal <a href="#">MagLev</a> AI training and inference infrastructure, involving collaborations across multiple groups, which resulted in two publications.</li> </ul>
May 2018 – Aug 2018	<b>NVIDIA, Santa Clara, USA</b> <i>Software Intern; <a href="#">Autonomous Vehicle Applied Research</a></i> <ul style="list-style-type: none"> <li>• <i>Mentors:</i> <a href="#">Dr. José M. Álvarez</a>, <a href="#">Dr. Adam Lesnikowski</a></li> <li>• <i>Role:</i> Research on approximating Bayesian Neural Networks for Active Learning which resulted in a publication, and was subsequently incorporated into the data annotation platform for the autonomous vehicles group.</li> </ul>

## Teaching

---

2019 – Now	<b>University of Tübingen, Germany</b> <i>Lead Teaching Assistant</i> <ul style="list-style-type: none"> <li>• <i>Apr 2023 – Jul 2023:</i> <a href="#">Autonomous Vision</a> (seminar, 5 teams of 2 students)</li> <li>• <i>Apr 2022 – Jul 2022:</i> <a href="#">Autonomous Vision</a> (seminar, 8 teams of 2 students)</li> </ul> <i>Teaching Assistant</i> <ul style="list-style-type: none"> <li>• <i>Apr 2021 – Jul 2021:</i> <a href="#">Computer Vision</a> (lecture, 150 students)</li> <li>• <i>Oct 2019 – Feb 2020:</i> <a href="#">Self-Driving Cars</a> (lecture, 80 students)</li> </ul>
------------	---

## Supervision

---

2019 – Now

### University of Tübingen, Germany

#### Master Thesis Advisor

- Apr 2025 – Now: [Micha Fauth](#) (Thesis: Waymo Open Scenario Generation Challenge)
- Nov 2024 – Now: [Long Nguyen](#) (Thesis: Waymo Open End-to-End Driving Challenge)
- Jul 2024 – Apr 2025: [Jens Beißwenger](#) (Thesis: Enhancing Model-Based Reinforcement Learning for Autonomous Driving)
- Jun 2024 – Feb 2025: [Melanie Schneider](#) (Thesis: [Generative Dataset Distillation: A New Hope?](#))
- Mar 2024 – Sep 2024: [Julian Zimmerlin](#) (Thesis: [Tackling CARLA Leaderboard 2.0 with End-to-End Imitation Learning](#))
- Feb 2023 – Aug 2023: [Daniel Dauner](#) (Thesis: [Vehicle Motion Planning using Data-Driven Simulation](#))
- Dec 2022 – Jun 2023: [Luis Winkelmann](#) (Thesis: [LiDAR-based Object Detection for Planning Transformers](#))
- Dec 2022 – Jun 2023: [Tim Schreier](#) (Thesis: [On Offline Evaluation of 3D Object Detection for Autonomous Driving](#))
- Nov 2022 – May 2023: [Siddharth Ramrakhiani](#) (Thesis: [Vision Transformers for Autonomous Driving](#))
- Nov 2022 – May 2023: [Jovan Cicvaric](#) (Thesis: [Generative Dataset Distillation](#))
- Mar 2021 – Sep 2021: [Bernhard Jaeger](#) (Thesis: [Expert Drivers for Autonomous Driving](#))
- Oct 2020 – Apr 2021: [Micha Schilling](#) (Thesis: [Visual Abstractions for Autonomous Driving](#))

## Academic Activities

---

### Workshop Organization

- ICCV 2025: [Learning to See: Advancing Spatial Understanding for Embodied Intelligence](#), 19.10.2025.
- CVPR 2025: [Embodied Intelligence for Autonomous Systems on the Horizon](#), 11.06.2025.
- CoRL 2024: [Safe and Robust Robot Learning for Operation in the Real World](#), 09.11.2024.
- ECCV 2024: [Autonomous Vehicles meet Multimodal Foundation Models](#), 29.09.2024.
- CVPR 2024: [Foundation Models for Autonomous Systems](#), 17.06.2024.
- CVPR 2023: [End-to-End Autonomous Driving: Emerging Tasks and Challenges](#), 18.06.2023.
- ICLR 2023: [Scene Representations for Autonomous Driving](#), 05.05.2023.

### Recorded Talks

- [Specializing General-Purpose Video Diffusion Models](#). ECCV Tutorial: [Recent Advances in Video Content Understanding and Generation](#), Milan, 30.09.2024.
- [Synthesizing Simulation Environments with Generative Models](#). CVPR Workshop on [Data-Driven Autonomous Driving Simulation](#), Seattle, 18.06.2024.
- [Benchmarking Foundation Models for Autonomous Driving](#). CVPR Tutorial: [Towards Building AGI in Autonomy and Robotics](#), Seattle, 18.06.2024.
- [Non-Reactive Autonomous Vehicle Simulation and Benchmarking](#). CVPR Workshop on [Autonomous Driving](#), Seattle, 17.06.2024.
- [Reading, Writing, and Reviewing for Robotics and Computer Vision Research](#). Sogang University Applied Data Engineering Seminar, Virtual, 07.06.2023.
- [End-to-End Driving with Attention](#). ICRA Workshop on [Scalable Autonomous Driving](#), London, 02.06.2023.
- [Imitation via Abstraction and Planning](#). ETH Computer Vision Lab, Zürich, 20.02.2023.

- [Imitation with Transformer-based Sensor Fusion for Autonomous Driving](#). University of Toronto AI in Robotics Seminar, Virtual, 28.03.2022.

## Reviewing and Service

- *Journal Reviewer*: T-PAMI, IJCV, T-RO, RA-L, T-IP, T-ITS, T-IV
- *Conference Reviewer*: CVPR, ICCV, ECCV, WACV, CoRL, ICRA, IROS, NeurIPS, ICLR, IV
- *Publicity Chair*: [RSS Pioneers 2024](#)
- *Program Chair*: [ICLR 2023 SR4AD Workshop](#), [ECCV 2024 MLLMAV Workshop](#)
- *Evaluator*: [ELLIS PhD Program, 2022-2024](#), [IMPRS-IS PhD Program, 2023-2024](#)

## Publications

---

- |      |  |
|------|--|
| 2025 | <ul style="list-style-type: none"> <li>[1] M. Fauth, L. Nguyen, B. Jaeger, D. Dauner, M. Igl, A. Geiger, and <b>K. Chitta</b>, “Shred: Synthesizing rule-based environments for driving,” in <i>Workshop on Autonomous Driving (WAD), Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2025.</li> <li>[2] L. Nguyen, M. Fauth, B. Jaeger, D. Dauner, M. Igl, A. Geiger, and <b>K. Chitta</b>, “Open x-av: Unifying end-to-end autonomous driving datasets,” in <i>Workshop on Autonomous Driving (WAD), Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2025.</li> <li>[3] C. Sima, <b>K. Chitta</b>, Z. Yu, S. Lan, P. Luo, A. Geiger, H. Li, and J. M. Alvarez, “Centaur: Robust end-to-end autonomous driving with test-time training,” in <i>Workshop on Test-time Scaling for Computer Vision, Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2025.</li> </ul>  |
| 2024 | <ul style="list-style-type: none"> <li>[4] L. Chen, P. Wu, <b>K. Chitta</b>, B. Jaeger, A. Geiger, and H. Li, “End-to-end autonomous driving: Challenges and frontiers,” <i>Transactions on Pattern Analysis and Machine Intelligence (T-PAMI)</i>, 2024.</li> <li>[5] <b>K. Chitta</b>, D. Dauner, and A. Geiger, “Sledge: Synthesizing driving environments with generative models and rule-based traffic,” in <i>European Conference on Computer Vision (ECCV)</i>, 2024.</li> <li>[6] D. Dauner, M. Hallgarten, T. Li, X. Weng, Z. Huang, Z. Yang, H. Li, I. Gilitschenski, B. Ivanovic, M. Pavone, A. Geiger, and <b>K. Chitta</b>, “Navsim: Data-driven non-reactive autonomous vehicle simulation and benchmarking,” in <i>Advances in Neural Information Processing Systems (NeurIPS)</i>, 2024.</li> <li>[7] S. Gao, J. Yang, L. Chen, <b>K. Chitta</b>, Y. Qiu, A. Geiger, J. Zhang, and H. Li, “Vista: A generalizable driving world model with high fidelity and versatile controllability,” in <i>Advances in Neural Information Processing Systems (NeurIPS)</i>, 2024.</li> <li>[8] M. Schneider, J. Cicvaric, A. Sauer, A. Geiger, and <b>K. Chitta</b>, “Generative dataset distillation: A new hope?” In <i>Workshop on the Dataset Distillation Challenge, European Conference on Computer Vision (ECCV)</i>, 2024.</li> <li>[9] C. Sima, K. Renz, <b>K. Chitta</b>, L. Chen, H. Zhang, C. Xie, J. Beißwenger, P. Luo, A. Geiger, and H. Li, “Drivelm: Driving with graph visual question answering,” in <i>European Conference on Computer Vision (ECCV)</i>, 2024.</li> <li>[10] J. Yang, S. Gao, Y. Qiu, L. Chen, T. Li, B. Dai, <b>K. Chitta</b>, P. Wu, J. Zeng, P. Luo, J. Zhang, A. Geiger, Y. Qiao, and H. Li, “Generalized predictive model for autonomous driving,” in <i>Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2024.</li> <li>[11] J. Zimmerlin, J. Beißwenger, B. Jaeger, A. Geiger, and <b>K. Chitta</b>, “Hidden biases of end-to-end driving datasets,” in <i>Workshop on Foundation Models for Autonomous Systems (FM4AS), Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2024.</li> </ul> |

2023	<p>[12] <b>K. Chitta</b>, A. Prakash, B. Jaeger, Z. Yu, K. Renz, and A. Geiger, "Transfuser: Imitation with transformer-based sensor fusion for autonomous driving," <i>Transactions on Pattern Analysis and Machine Intelligence (T-PAMI)</i>, 2023.</p> <p>[13] D. Dauner, M. Hallgarten, A. Geiger, and <b>K. Chitta</b>, "Parting with misconceptions about learning-based vehicle motion planning," in <i>Conference on Robot Learning (CoRL)</i>, 2023.</p> <p>[14] B. Jaeger, <b>K. Chitta</b>, and A. Geiger, "Hidden biases of end-to-end driving models," in <i>International Conference on Computer Vision (ICCV)</i>, 2023.</p> <p>[15] T. Schreier, K. Renz, A. Geiger, and <b>K. Chitta</b>, "On offline evaluation of 3d object detection for autonomous driving," in <i>Workshop on Robustness and Reliability of Autonomous Vehicles in the Open-world (BRAVO), International Conference on Computer Vision (ICCV)</i>, 2023.</p>
2022	<p>[16] N. Hanselmann, K. Renz, <b>K. Chitta</b>, A. Bhattacharyya, and A. Geiger, "King: Generating safety-critical driving scenarios for robust imitation via kinematics gradients," in <i>European Conference on Computer Vision (ECCV)</i>, 2022.</p> <p>[17] K. Renz, <b>K. Chitta</b>, O.-B. Mercea, A. S. Koepke, Z. Akata, and A. Geiger, "Plant: Explainable planning transformers via object-level representations," in <i>Conference on Robot Learning (CoRL)</i>, 2022.</p>
2021	<p>[18] <b>K. Chitta</b>, J. M. Alvarez, E. Haussmann, and C. Farabet, "Training data subset search with ensemble active learning," <i>Transactions on Intelligent Transportation Systems (T-ITS)</i>, 2021.</p> <p>[19] <b>K. Chitta</b>, A. Prakash, and A. Geiger, "Neat: Neural attention fields for end-to-end autonomous driving," in <i>International Conference on Computer Vision (ICCV)</i>, 2021.</p> <p>[20] A. Prakash, <b>K. Chitta</b>, and A. Geiger, "Multi-modal fusion transformer for end-to-end autonomous driving," in <i>Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2021.</p> <p>[21] A. Sauer, <b>K. Chitta</b>, J. Muller, and A. Geiger, "Projected gans converge faster," in <i>Advances in Neural Information Processing Systems (NeurIPS)</i>, 2021.</p> <p>[22] M. A. Weis, <b>K. Chitta</b>, Y. Sharma, W. Brendel, M. Bethge, A. Geiger, and A. S. Ecker, "Benchmarking unsupervised object representations for video sequences," <i>Journal of Machine Learning Research (JMLR)</i>, 2021.</p>
2020	<p>[23] A. Behl, <b>K. Chitta</b>, A. Prakash, E. Ohn-Bar, and A. Geiger, "Label efficient visual abstractions for autonomous driving," in <i>International Conference on Intelligent Robots and Systems (IROS)</i>, 2020.</p> <p>[24] <b>K. Chitta</b>, J. M. Alvarez, and M. Hebert, "Quadtree generating networks: Efficient hierarchical scene parsing with sparse convolutions," in <i>Winter Conference on Applications of Computer Vision (WACV)</i>, 2020.</p> <p>[25] E. Haussmann, M. Fenzi, <b>K. Chitta</b>, J. Ivaneky, H. Xu, D. Roy, A. Mittel, N. Koumchatzky, C. Farabet, and J. M. Alvarez, "Scalable active learning for object detection," in <i>Intelligent Vehicles Symposium (IV)</i>, 2020.</p> <p>[26] E. Ohn-Bar, A. Prakash, A. Behl, <b>K. Chitta</b>, and A. Geiger, "Learning situational driving," in <i>Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2020.</p> <p>[27] A. Prakash, A. Behl, E. Ohn-Bar, <b>K. Chitta</b>, and A. Geiger, "Exploring data aggregation in policy learning for vision-based urban autonomous driving," in <i>Conference on Computer Vision and Pattern Recognition (CVPR)</i>, 2020.</p>

- |      |  |
|------|--|
| 2018 | <p>[28] <b>K. Chitta</b>, “Targeted kernel networks: Faster convolutions with attentive regularization,” in <i>Workshop on Compact and Efficient Feature Representation and Learning in Computer Vision (CEFRL), European Conference on Computer Vision (ECCV)</i>, 2018.</p> <p>[29] <b>K. Chitta</b>, J. M. Alvarez, and A. Lesnikowski, “Deep probabilistic ensembles: Approximate variational inference through kl regularization,” in <i>Workshop on Bayesian Deep Learning (BDL), Conference on Neural Information Processing Systems (NeurIPS)</i>, 2018.</p> |
| 2016 | <p>[30] <b>K. Chitta</b> and N. N. Sajjan, “A reduced region of interest based approach for facial expression recognition from static images,” in <i>IEEE Region-10 Conference (TENCON)</i>, 2016.</p>   |

## References

---

<b>Prof. Andreas Geiger</b> .	Head of the Dept. of Computer Science, University of Tübingen.	a.geiger@uni-tuebingen.de
<b>Prof. Marco Pavone</b> .	Director, Autonomous Vehicles Research, NVIDIA.	mpavone@nvidia.com
<b>Dr. José M. Álvarez</b> .	Director, Autonomous Vehicle Applied Research, NVIDIA.	josea@nvidia.com
<b>Prof. Hongyang Li</b> .	Assistant Professor, University of Hong Kong.	hongyang@hku.hk